NEATHACK · CFO OPERATING PLAN

PriorProof

Prove or refuse.
Then earn the right to scale.

US specialty-practice evidence packets · USD
October 2026–September 2027
Prepared October 12, 2026

$106k modeled first-year revenue · 60% gross margin
$48k base funding with reserve · $73k downside

Independent pressure test. No booked revenue, compliance certification, or approval guarantee. Public facts and planning assumptions are separated throughout.

1. Executive recommendation

PriorProof should launch as an evidence-checking assistant for one specialty and a small, explicit set of payer policies—not as an autonomous approval or appeal engine. Start with complex musculoskeletal procedure authorizations, which align more closely with the policy example on the live site, unless founder interviews establish a better first market (recommendation). Sell auditable documentation and named evidence gaps; do not promise authorization, reimbursement, or appeal success.

The public site describes a synthetic-chart demo and labels its sample “Canned illustration. No live engine call.” That supports the product concept, not a demonstrated production engine, performance benchmark, clinical-safety claim, or measured packet cost [1].

The recommended base case is $105,545 revenue, 15,396 billable packets, $42,305 COGS, $63,240 gross profit, and 59.9% gross margin in October 2026–September 2027 (assumptions; packet count rounded). Monthly operating surplus turns positive in April 2027, month 7; net cash turns positive in May, month 8 (assumptions). The twelve-month net cash movement is negative $10,927 (assumption).

Opening cash is unknown. With $50,000 unrestricted starting cash (placeholder assumption), the base case has a $12,228 April trough and ends September at $39,073 (assumptions). Required opening funds are $37,772 before a reserve, or $47,772 including a $10,000 reserve (assumptions). Budget approximately $75,000 to cover the downside plus that reserve (recommendation based on assumptions), rather than calling this a $1,500-buffer business.

All commercial prices, customer counts, adoption, costs except published token-rate references, cash timing, budgets, and projections below are assumptions. Calculated outputs inherit that label. They are not agreed customer commitments, booked results, market quotes, or certified compliance costs.

2. What the draft gets right—and what must change

The problem is real. An AMA article dated April 17, 2025 reports 39 prior authorizations per physician per week, 13 hours of physician/staff time, and 40% of physicians with staff working exclusively on prior authorization [2]. Cite that dated article and its survey context; do not label these numbers as the latest 2025 survey results or turn a per-physician measure into a practice-level sales forecast.

The 2024 CAQH Index reports medical-provider prior-authorization transaction cost of $12.88 manual, $8.93 partially electronic, and $5.38 fully electronic; the provider opportunity from fully manual to fully electronic is $7.50 per transaction. See the report’s Average Cost and Savings Opportunity per Transaction table, printed page 51 (PDF page 56) [3]. Replace $10.91/$2.17 in the draft with this consistent primary-source pair. These are industry transaction benchmarks, not demonstrated savings from PriorProof and not a floor for every specialty case. CAQH explicitly excludes information gathering, follow-up, and system costs from this transaction-cost table [3]. That makes it particularly unsuitable as an estimate of the whole evidence-assembly workflow.

KFF’s 2022 Medicare Advantage analysis verifies the 7.4% denial rate, 9.9% appeal rate among denials, and 83.2% overturn rate among appeals [4]. The last number is conditional on appeals actually selected and filed. It does not imply 83.2% of unappealed cases would win, or that PriorProof will win four in five appeals. Do not include recovered reimbursement in this forecast.

HHS OIG’s 13% finding concerned sampled prior-authorization denials by 15 large Medicare Advantage organizations during June 1–7, 2019; the report was published in 2022 [5]. It is evidence of a historical access problem, not a current universal denial-error rate.

The draft’s arithmetic is broadly correct: its ramp totals 21,650 packets, $194,850 revenue, $181,860 gross profit, and $142,260 net cash after its stated founder draw and infrastructure (founder assumptions, reference only). The issue is unsupported inputs, not multiplication. The original $9 realized price, $0.60 COGS, and first-month 200 paid packets are not adopted as measured facts.

Twenty model calls do not determine cost. Tokens per call, chart length, retries, escalation, extraction, retention controls, and human work do. Two model passes may share errors; agreement is not statistical independence. A real quotation can still be the wrong patient, outdated, negated, inadequate in dose/duration, or insufficient for the criterion.

“Submission-ready” is too strong for a criteria-only matrix. Patient and member identifiers, plan effective dates, codes, ordering clinician details, forms, benefit rules, network restrictions, deadlines, and reviewer sign-off may still be required. Call the initial product an evidence packet for staff review. A gap can represent care or testing that has not happened; the tool must request genuine evidence, not suggest that a fabricated sentence would close it.

3. Product, buyer, and route to market

The initial job is narrow: retrieve the applicable public payer policy, pin its version and effective date, extract the criteria, compare each against the right chart, cite the supporting span, return genuine gaps, and export an auditable draft for practice staff. A named gap is a useful outcome. A red matrix is not a technical failure and must never be pressured into green to earn a fee.

Begin with specialty practice administrators and prior-auth leads who can run a limited pilot without an enterprise integration project. RCM companies are the second channel: one agreement can distribute the product, but its BAA posture does not cover PriorProof, and a forty-practice customer does not imply forty immediate deployments. Model actual routing volume, not the buyer’s total footprint.

Working launch sequence: three readiness months followed by paid processing in January 2027 (assumption, not an announced launch). Seek two new Pro accounts per month in January–March, three per month in April–June, and four per month in July–September; 27 gross new Pro accounts before churn (assumptions). At a 20% qualified-opportunity-to-paid conversion rate, that requires 10, 15, then 20 genuinely qualified clinic opportunities per month respectively (sales-target assumptions; not a verified pipeline). Seek one RCM contract starting April and one starting July (assumptions).

The base case assumes 2% monthly Pro logo churn, 0.5% Scale churn, 100 monthly attempted runs per Pro account, and 600 per Scale account (assumptions). It models 90% technical completion as billable (assumption). Expected retained accounts can be fractional; real contracts cannot. By month 12 the model represents about 25.27 retained Pro accounts and 1.97 Scale accounts, about 3,436 billable packets, and $23,465 monthly revenue (assumptions). Pro and Scale dominate revenue; PAYG contributes only $6,480 during the year (assumption).

Existing staff, payer portals, electronic transaction tools, and broader automation are alternatives. Infinitus, for example, advertises prior-auth call and follow-up automation [11]. PriorProof's differentiation is evidence traceability and conservative refusal, not a claim that it automates the entire workflow or has no competitors. CMS's 2024 final rule generally phases relevant payer APIs in beginning January 2027 and excludes drugs from its prior-auth provisions [12]. Treat interoperable submission as a potential partner channel and competitive pressure, not as a mandate to buy PriorProof or a rule applying identically to all commercial and drug workflows.

4. Recommended pricing and billable definition

PAYG: $12 per usable evidence packet or named-gap matrix; no commitment (assumption). A technical crash, missing output, or unusable export is not billable. Include a clear policy for corrections and repeated submissions so retries are not sold as new packets.

Pro: retain $349 per month including 50 billable packets, then $7 per additional packet (assumption). At 100 attempts and 90% billable completion, the account produces 90 billable outputs and $629 monthly revenue, equivalent to $6.99 per output (assumptions). The subscription minimum is payable even when usage is lower, subject to the service contract.

Scale: start at $6 per billable packet with a $3,000 monthly commitment, equivalent to 500 included-at-rate packets, not a $5–$7 promise with no minimum (assumption). The modeled 600 attempts produce 540 billable outputs and $3,240 monthly revenue per active account (assumptions). Charge custom onboarding and integrations separately only after defining scope; no such fees or projects are assumed in year one.

The year-one realized blend is $6.86, not $9 (assumption), because most output comes through Pro and Scale. Holding volume fixed, ignoring tier mix is enough to overstate revenue. The model calculates each tier's revenue independently before summing it.

All real-PHI tiers require an appropriate customer BAA, eligible covered subprocessors, security controls, and operating procedures before any real-chart processing. A BAA is not a Scale-only feature [6–9]. No-PHI sandbox access is a separate demo experience, not a billable patient submission product.

Pricing is a willingness-to-pay hypothesis. At an assumed $40/hour practice labor cost, a $6.99 output needs about 10.5 minutes of net labor savings just to pay its fee; a $12 PAYG output needs 18 minutes (assumptions). Net savings must subtract staff verification, upload, corrections, and submission. Measure that in pilots. The draft's $90 complex-case cost is an unmeasured estimate, not a defensible market ceiling. Do not use it as proven ROI.

5. Packet COGS: measure tokens, then include delivery

Reference model: published GPT-4.1 mini pricing is $0.40 per million input tokens and $1.60 per million output tokens [10]. This is a rate reference, not an assertion about the deployed engine or its suitability. The forecast assumes 180,000 aggregate input tokens and 24,000 aggregate output tokens across a run, no caching discount, and a 1.2 retry multiplier (assumptions). Base LLM cost is $0.1104 before retries and $0.1325 after retries per attempted run (assumptions).

Per attempted run, add extraction/OCR $0.08, variable hosting/storage $0.04, support $0.25, and paid coordinator review averaging two minutes at $45/hour, or $1.50 (assumptions). Total delivery cost is $2.0025 per attempt (assumption). Review is incremental external delivery labor; do not count the founder as both this paid reviewer and a free resource. Coordinator QA does not replace the practice clinician's clinical responsibility.

At 90% billable completion, delivery cost becomes $2.225 per billable packet (assumption). This denominator pays for failed attempts as well as completed outputs. Add a blended payment fee assumption of 2.5% of revenue plus $0.03 per billable packet; this is an amortized payment-cost allowance, not a processor quote. Variable COGS averages approximately $2.43 per billable output in the base case (assumption).

Fixed serving infrastructure is $150/month during readiness and $500/month after launch, classified above the gross-profit line (assumptions). That includes a provisional allowance for secure hosting, observability, and the HIPAA-capable platform posture; obtain quotes before claiming adequacy. Annual allocated COGS is $2.75 per packet and annual gross margin 59.9%; month-12 allocated COGS is $2.57 and gross margin 62.3% (assumptions). These are different allocation periods, not inconsistent margins.

Human effort is the biggest risk. At four average review minutes, delivery cost before payment fees rises to $3.89 per billable output; at eight minutes it is $7.22, already above the Scale price before payment fees or fixed infrastructure (assumptions). At month-12 base volume, paid review is about 127 hours/month (assumption). A solo founder cannot silently supply this work while also building and selling. Reduce the average only after validation supports it; otherwise reprice or narrow the use case.

The $0.60 draft estimate may be plausible for raw software-only consumption on some workloads. It is not an all-in commercial delivery COGS estimate. Logging real aggregate tokens, failed runs, review time, support contacts, and costs per usable packet is required to replace this provisional model.

6. Operating budgets and break-even

Cash operating budgets during readiness are founder allowance $3,000, selling $600, compliance/insurance $400, and backoffice $200 per month, totaling $4,200 (assumptions), plus $150 fixed serving COGS. Production budgets are founder $3,000, selling $1,000, compliance/insurance $600, and backoffice $200, totaling $4,800 (assumptions), plus $500 serving COGS. The steady fixed monthly cash burden is therefore $5,300 (assumption).

A separate $12,000 readiness budget is paid $6,000, $4,000, then $2,000 in the first three months (assumptions). Suggested scope: $3,000 legal/contracts, $5,000 security assessment/remediation, and $4,000 independent workflow/evidence evaluation (assumptions, not vendor quotes). Founder-led engineering is funded through the modest founder allowance; no market-rate engineering employee or custom EHR integration build is included. Broad enterprise integrations or a requirement for a costly platform plan would need additional funding.

Using the month-12 mix, realized price is $6.83, variable cost $2.43, and contribution $4.40 per billable packet (assumptions). Serving-only break-even is approximately 114 billable packets/month; serving plus founder allowance is about 795; the full modeled operating burden requires approximately 1,204 (assumptions). These are constant-mix, constant-usage equivalents, not an exact threshold for every Pro minimum or Scale contract composition. Startup costs and collection lags do not disappear at operating break-even.

The founder allowance is included as a cash burden for viability. Depending on the legal entity, an owner distribution is not a P&L expense; a corporate salary may incur additional taxes and benefits. This is a cash-planning operating surplus, not a GAAP profit claim. Company income taxes, payroll gross-up, financing, and major capital expenditure are excluded and must be added if applicable.

7. Twelve-month forecast and runway

The detailed forecast is saved on the business-plan page and included below in this document. Months one through three are readiness with no paid packet revenue (assumption). PAYG begins at 20 billable outputs in January and adds 10/month (assumptions). The clinic and RCM cohort records drive all other packet volume; no packet ramp is typed directly into the adopted forecast.

Clinic and PAYG cash arrives in the earned month, while Scale invoices collect the following month (assumptions). September ends with $6,367 Scale receivables, explaining the gap between year-one operating surplus of negative $4,560 and net cash movement of negative $10,927 (assumptions). Collections, delivery costs, startup expense, and founder cash allowance all feed the cash balance.

Base case: $105,545 year-one revenue, $23,465 month-12 revenue, $12,228 minimum cash, and $39,073 ending cash on a $50,000 starting balance (assumptions). No zero-cash crossing appears through September 2027. This supports twelve covered months only, not an infinite runway claim. Monthly operating surplus becomes positive in April and cash flow in May; cumulative cash movement remains negative throughout the horizon (assumptions).

Downside: shift launch and every prospective cohort three months later, halve acquisition, and double average review time to four minutes (assumptions). Year-one revenue falls to $21,095, allocated COGS to $5.38/packet, and gross margin to 22.6% (assumptions). With $50,000 starting cash, the $10,000 reserve is breached in April and cash goes below zero in June, month 9; September ending cash is negative $12,841 (assumptions). Negative balances show an unfunded plan, not permission to keep spending. It needs $62,841 opening cash before a reserve, or $72,841 including the reserve (assumptions).

Upside: retain the base launch timing, increase expected acquisition 50%, and reduce average paid review to one minute through validated process improvements (assumptions). Revenue is $158,318, margin 73.6%, and ending cash $89,228 on the same $50,000 opening balance (assumptions). Net cash first becomes positive in April. The downside and upside are expectations, not literal half or one-and-a-half enterprise contracts. Actual deals remain discrete, so individual sale timing can cause a sharper cash swing.

Zero-sales fallback: if the entire year remained in the lower readiness cost posture, $12,000 setup plus twelve months at $4,350 requires $64,200 before a reserve (assumption). Do not keep spending simply because the optimistic packet ramp once said cash flow would turn positive.

Base-case monthly forecast

Every cell below is a modeled assumption or calculated result of assumptions. USD. Packets rounded; calculations retain fractional expected volumes. Founder allowance and readiness costs included. Opening cash $50,000 (placeholder assumption).

Month	Packets	Revenue	Gross profit	Net cash	Closing cash
Oct 2026	0	$0	−$150	−$10,350	$39,650
Nov	0	$0	−$150	−$8,350	$31,300
Dec	0	$0	−$150	−$6,350	$24,950
Jan 2027	200	$1,498	$510	−$4,290	$20,660
Feb	386	$2,851	$1,408	−$3,392	$17,268
Mar	569	$4,179	$2,291	−$2,509	$14,759
Apr	1,379	$9,352	$5,509	−$2,531	$12,228
May	1,640	$11,233	$6,753	$1,969	$14,197
Jun	1,897	$13,078	$7,974	$3,190	$17,387
Jul	2,778	$18,757	$11,523	$3,500	$20,887
Aug	3,110	$21,134	$13,092	$8,324	$29,211
Sep	3,436	$23,465	$14,630	$9,862	$39,073

Gross profit is revenue minus variable and fixed serving COGS. Net cash additionally deducts cash operating budgets and startup readiness spending, and adjusts for Scale collections.

Scenario comparison

All figures are assumptions. All cases start with $50,000 cash. Negative ending cash is an unfunded spending plan.

Scenario	Year revenue	Gross margin	Ending cash	Funds + reserve
Base	$105,545	59.9%	$39,073	$47,772
Downside	$21,095	22.6%	−$12,841	$72,841
Upside	$158,318	73.6%	$89,228	$40,942

8. What must be true before this forecast becomes investable

HIPAA is an operating program, not a signature. HHS requires risk analysis and appropriate safeguards; de-identification requires Safe Harbor or Expert Determination, not merely deleting a patient name [6–8]. An application that receives PHI to de-identify it is still handling PHI at that point. Map every PHI path, including storage, OCR, queues, backups, analytics, support, and logs.

Vercel currently describes a paid HIPAA add-on for Pro and Enterprise BAA support; using it does not itself make the application compliant [9]. OpenAI describes BAA support for eligible API organizations, with covered-service and configuration requirements [6]. Confirm the signed terms, covered endpoints, retention configuration, and the actual service architecture rather than assuming an ordinary account is sufficient. Obtain customer and subcontractor agreements, restrict access, secure transport and storage, control PHI in logs, define retention/deletion, handle incidents, and complete security review with appropriate professional advice.

Use the synthetic demo for neatHack. Before production, run an independently adjudicated evaluation over 100 policy/chart case pairs in the first narrow workflow (evaluation-size assumption), including negation, stale policies, wrong patient/time, contradictory notes, numeric thresholds, dose/duration, missing evidence, and technical failure. Report every observed unsupported green and missed gap; observing none in a finite sample is not proof of zero population error. Public performance claims must reflect the test design and observed denominators.

Proposed commercial gates are 90% usable technical completion, net practice staff time savings sufficient to exceed the actual tier fee, no known unresolved unsupported-green failures in the release evaluation, acceptable buyer review burden, and paid pilot renewal (targets/assumptions, not achieved results). Record reviewer corrections, refusal rates by cause, repeat submissions, net staff minutes, token cost, and support work. Validate payer-policy applicability and the billable distinction between a useful refusal and a broken run.

Earn repeat paid use in the narrow specialty before expanding the policy library. Version policies, save provenance, require staff sign-off, and exclude autonomous treatment advice or guaranteed coverage. Add appeals only after a denial-specific workflow is validated; the launch plan books neither appeal-success bonuses nor recovered practice revenue.

My decision: proceed with a tightly scoped paid-pilot product, conditional on the above gates. Keep the $12/$349/$6 commercial tests, secure $50,000 for the base or approximately $75,000 for a downside-resilient launch with reserve (recommendations using assumptions), and reject the 93%-margin, almost-no-capital narrative until delivery and buyer evidence supports it. The most valuable next step is an instrumented pilot that tests willingness to pay and review burden simultaneously—not another unsourced market-size slide.

Sources

[1] PriorProof live product page; retrieved October 12, 2026
https://priorproof.vercel.app

[2] AMA, Don’t fall for these myths on prior authorization, April 17, 2025
https://www.ama-assn.org/practice-management/prior-authorization/don-t-fall-these-myths-prior-authorization

[3] 2024 CAQH Index, medical cost table, printed page 51 (PDF page 56)
https://www.caqh.org/hubfs/Index/2024%20Index%20Report/CAQH_IndexReport_2024_FINAL.pdf

[4] KFF, Medicare Advantage denials in 2022, August 8, 2024
https://www.kff.org/medicare/medicare-advantage-plans-denied-a-larger-share-of-prior-authorization-requests-in-2022-than-in-prior-years/

[5] HHS OIG, OEI-09-18-00260, April 27, 2022; sampled 2019 denials
https://oig.hhs.gov/reports/all/2022/some-medicare-advantage-organization-denials-of-prior-authorization-requests-raise-concerns-about-beneficiary-access-to-medically-necessary-care/

[6] OpenAI, Getting a BAA for the OpenAI API; current provider guidance
https://help.openai.com/en/articles/8660679-getting-a-business-associate-agreement-for-the-openai-api

[7] HHS, Guidance on Risk Analysis
https://www.hhs.gov/hipaa/for-professionals/security/guidance/guidance-risk-analysis/index.html

[8] HHS, Guidance Regarding Methods for De-identification
https://www.hhs.gov/hipaa/for-professionals/special-topics/de-identification/index.html

[9] Vercel, HIPAA compliance guide; current provider guidance
https://vercel.com/kb/guide/hipaa-compliance-guide-vercel

[10] OpenAI, GPT-4.1 mini published token pricing
https://developers.openai.com/api/docs/models/gpt-4.1-mini

[11] Infinitus, Prior authorization follow-up automation; vendor-described scope
https://www.infinitus.ai/solutions/prior-authorization/

[12] CMS, CMS-0057-F fact sheet, 2024 final rule scope and dates
https://www.cms.gov/files/document/fact-sheet-cms-interoperability-and-prior-authorization-final-rule-cms-0057-f.pdf